

## CS224 – Spring 2021 – Lab #6

### Examining the Effect of Cache Parameters and Program Factors on Cache Hit Rate

#### Dates:

- Section 1: Mon, 3 May, 8:30-12:20 in EA-Z04
- Section 2: Wed, 28 Apr, 13:30-17:20 in EA-Z04
- Section 3: Tue, 27 Apr, 13:30-17:20 in EA-Z04
- Section 4: Mon, 3 May, 13:30-17:20 in EA-Z04
- Section 5: Fri, 30 Apr, 8:30-12:20 in EA-Z04
- Section 6: Fri, 30 Apr, 13:30-17:20 in EA-Z04

**Prelim Submission Date:** April 27, Tuesday, 13:30(See the submission instructions below)

**Physical Lab Administration:** No physical lab: All labs are online. Your TAs will be available during your lab hours. Other times you may reach them by email.

#### TAs:

- Section 1: Zülal Bingöl, Ergün Batuhan Kaynak
- Section 2: Zülal Bingöl, Kenan Çağrı Hırlak
- Section 3: Alper Şahistan, Kenan Çağrı Hırlak
- Section 4: Ege Berkay Gülcan, Ergün Batuhan Kaynak
- Section 5: Hüseyin Eren Çalık, Yusuf Dalva
- Section 6: Ziya Erkoç, Yusuf Dalva

#### TA name (x No of labs) email address:

- Alper Şahistan (x1): alper.sahistan@bilkent.edu.tr
- Ege Berkay Gülcan (x1): berkay.gulcan@bilkent.edu.tr
- Ergün Batuhan Kaynak (x2): batuhan.kaynak@bilkent.edu.tr
- Hüseyin Eren Çalık (x1): eren.calik@bilkent.edu.tr
- Kenan Çağrı Hırlak (x2): cagri.hirlak@bilkent.edu.tr
- Yusuf Dalva (x2): yusuf.dalva@bilkent.edu.tr
- Ziya Erkoç (x1): ziya.erkoc@bilkent.edu.tr
- Zülal Bingöl (x2): zulal.bingol@bilkent.edu.tr

**Purpose:** Studying the effect of various cache design parameters. The first part includes problem solving and writing a program. The second part involves execution of the program and preparing a report.

In your solutions and report make sure that you have proper tables, page numbering and understandable explanation with good writing. All tables must have a subtitle and table number. In tables columns must have meaning names/headers.

## Summary

**Prelim (Part 1)** (50 points): Involves problem solving related to cache memory design and a program written for cache testing.

**Lab Work (Part 2)** (50 points): Experiments with caching and experiment report.

**PRELIM DUE DATE: SAME FOR ALL SECTIONS:** As given above.

## Submit Problem Solutions and Your Code for MOSSsimilarity testing

**Report:** Put your solutions for Part 1.1, 1.2, 1.3 together into a PDF file. Use filename

**StudentID\_FirstName\_LastName\_SecNo\_Lab6\_prelim\_report.pdf** [pdf FILE as its extension suggests, which contains all the work done for the problem solutions Part].

**Code:** Put your MIPS code for Part 1.4 into a .txt file. Use filename

**StudentID\_FirstName\_LastName\_SecNo\_Lab6\_prelim\_code.txt** [ANOTEPAD FILE as its extension suggests, which contains the Program Code Part]

**LAB DUE DATE: VARIES FROM SECTION TO SECTION.**

## Submit Experiment Report and Your Code for MOSSsimilarity testing

**Report:** Put your experiment results for Part 2 (including the tables and graphs) in a single PDF file.

Use filename **StudentID\_FirstName\_LastName\_SecNo\_Lab6\_lab\_report.pdf** [pdf FILE as its extension suggests, which contains all the work done for the Lab Experiment Report Part].

**Code:** Put your MIPS code for Part 1.4 into a .txt file. Use filename

**StudentID\_FirstName\_LastName\_SecNo\_Lab6\_lab\_code.txt** [ANOTEPAD FILE as its extension suggests, which contains the Program Code Part]

Your program (code) will be compared against all the other programs in the class, by the MOSS program, to determine how similar it is (as an indication of plagiarism). So be sure that the code you submit is code that you actually wrote yourself! The same type of comparison is also planned for the reports.

## PRELIM (50 points)

### Part 1. Cache Memory Problems and Program (50 points)

You have to provide a neat presentation prepared by Word or a word processor with similar output quality. Handwritten answers will not be accepted. At the top of the paper on left provide the following information. Please make sure that this info is there for proper grading of your work, otherwise some points will be taken off.

CS224

Section No.: ...

Spring 2021

Lab No.:

Your Full Name/Bilkent ID:

**1. ( 5 points: With 3 or more errors you get 0 points. Otherwise full point.)** Fill in the empty cells of the following table. Assume that main memory size is 0.5 GB. **Index Size:** No. of bits needed to express the set number in an address, **Block Offset:** No. of bits needed to indicate the word offset in a block, **Byte Offset:** No. of bits needed to indicate the byte offset in a word. **Block Replacement Policy Needed:** Indicate if a block replacement policy such as FIFO, LRU, LFU (Least Frequently Used) etc. is needed (yes) or not (no). If some combinations are not possible mark them.

No.	Cache Size KB	N way cache	Word Size in bits	Block size (no. of words)	No. of Sets	Tag Size in bits	Index Size (Set No.) in bits	Word Block Offset Size in bits <sup>1</sup>	Byte Offset Size in bits <sup>2</sup>	Block Replacement Policy Needed (Yes/No)
1	8	1	8	8	2 <sup>10</sup>	16	10	3	0	No
2	8	2	16	8	2 <sup>8</sup>	17	8	3	1	Yes
3	8	4	16	4	2 <sup>6</sup>	20	6	2	1	Yes
4	8	Full	16	4	2 <sup>0</sup>	26	0	2	1	Yes
9	32	1	16	2	2 <sup>13</sup>	14	13	1	1	No
10	32	2	16	2	2 <sup>12</sup>	15	12	1	1	Yes
11	32	4	8	8	2 <sup>8</sup>	18	8	3	0	Yes
12	32	Full	8	8	2 <sup>0</sup>	26	0	3	0	Yes

<sup>1</sup> Word Block Offset Size in bits:  $\log_2(\text{No. of words in a block})$

<sup>2</sup> Byte Offset Size in bits:  $\log_2(\text{No. of bytes in a word})$

**2. (5 points: With 3 or more errors you get 0 points. Otherwise full point.)** Consider the following MIPS code segment. (Remember MIPS memory size is 4 GB.) Cache capacity is 8 words, Block size: 4 words, N= 1.

```
loop:    addi    $t0, $0, 5
        beq     $t0, $0, done
        lw      $t1, 0xA4($0)
        lw      $t2, 0xA8($0)
        lw      $t3, 0xAC($0)
        addi    $t0, $t0, -1
        j       loop
done:
```

- a. In the following table indicate the type of miss, if any: Compulsory, Conflict, Capacity.

Instruction	Iteration No.				
	1	2	3	4	5
lw \$t1, 0xA4(\$0)	Compulsory				
lw \$t2, 0xA8(\$0)					
lw \$t3, 0xAC(\$0)					

- b. What is the memory size of one set in number of bits? What is the total cache memory (SRAM) size in number of bits? Note: Include the V bit in your calculations. Show the details of your calculation.

- c. State the number of AND and OR gates, EQUALITY COMPARATORS and MULTIPLEXERS needed to implement the cache memory. No drawing is needed.

3. (5 points: With 3 or more errors you get 0 points. Otherwise full point.) Consider the above MIPS code segment. The cache capacity is 2 words, block size is 1 word. N= 2. The block replacement policy is LRU.

- a. In the following table indicate the type of miss, if any: Compulsory, Conflict, Capacity.

Instruction	Iteration No.				
	1	2	3	4	5
lw \$t1, 0xA4(\$0)					
lw \$t2, 0xA8(\$0)					
lw \$t3, 0xAC(\$0)					

- b. How many bits are needed for the implementation of LRU policy: for a set, for the entire cache memory? What is the total cache memory size in number of bits? Include the V bit and the bit(s) used for LRU in your calculations. Show the details of your calculation.

- c. State the number of AND and OR gates, EQUALITY COMPARATORS and MULTIPLEXERS needed to implement the cache memory. No drawing is needed.

4. (35 points) Write a program to find the average of the elements of a square matrix. Provide a user interface for user interaction to demonstrate that your program is working properly. Assume that in the main memory matrix elements are placed **row-by-row**. Create an array for the matrix elements and initialize them row by row with consecutive values. For example, a 3 by 3 (N= 3) matrix would have the following values.

1	2	3
4	5	6
7	8	9

The row by row placement means that you will have the values of the above 3 x 3 matrix are stored as follows in the memory.

<b>Matrix Index (Row No., Col. No.)</b>	(1, 1)	(1, 2)	(1, 3)	(2, 1)	(2, 2)	(2, 3)	(3, 1)	(3, 2)	(3, 3)
<b>Displacement With respect the beginning of the array containing the matrix</b>	0	4	8	12	16	20	24	28	32
<b>Value stored</b>	1	2	3	4	5	6	7	8	9

In this configuration accessing the matrix element (i, j), where i is the row index and j is the column index, simply involves computation of its displacement from the beginning of the array that stores the matrix elements. For example, the displacement of the matrix element with the index (i, j) with respect to the beginning of the array is  $(i - 1) \times N \times 4 + (j - 1) \times 4$ , for a matrix of size  $N \times N$ .

Your user interface must provide at least the following functionalities,

1. Ask the user the matrix size in terms of its dimensions (N),
2. Allocate an array with proper size using syscall code 9,
3. Display desired elements of the matrix by specifying its row and column member
4. Obtain the average (arithmetic mean) of the matrix elements in a row-major (row by row) fashion, and display. That is, you should traverse the matrix row-by-row.
5. Obtain the average (arithmetic mean) of matrix elements in a column-major (column by column) fashion, and display. That is, you should traverse the matrix column-by-column.

## LAB WORK (50 points)

### Part 2. [50 pts] Experiments with Data Cache Parameters

Run your program with two reasonably large different matrix sizes that would provide meaningful observations. Modify your original program if needed. For large matrix initialization you may use a loop rather than user interface.

#### Report for Matrix Size 1: 25 Points

#### Report for Matrix Size 2: 25 Points

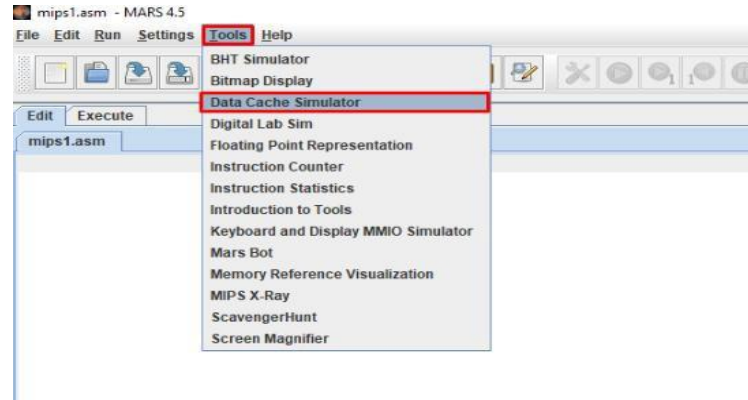
As described above make sure that you have an easy to follow presentation with numbered tables having proper heading etc.

Make sure that you find the average (arithmetic mean) of matrix elements by performing row-major and column-major traversal of the matrix. Note that the row-major averaging involves a simple array traversal from the beginning to the end; however, the column-major averaging is somewhat tricky.


- a) **Direct Mapped Caches:** For the matrix sizes you have chosen, conduct tests with various cache sizes and block sizes, to determine the hit rate, miss rate and number of misses. Use at least 5 different cache sizes and 5 different block sizes (make sure your values are reasonable) in order to obtain curves like those of Figure 8.18 (see below) in the textbook. This figure shows that as we increase

the block size the number of compulsory misses decrease and it decreases the overall miss rate. If we continue in that fashion after a point number of conflict misses start to increase and the overall miss rate starts to increase. Make a 5 x 5 table with your values, with miss rate and # of misses as the data at each row-column location. Make the graph of miss rate versus block size, parameterized by cache size, like Figure 8.18 both for row-major and column-major averaging.

Hint: You can reach the Cache Simulator from MARS/Tools/Data Cache Simulator as shown in the following image:



- b) **Fully Associative Caches:** Pick 3 of your parameter points obtained in part a) for **column-major averaging**, one with good hit rate, one with medium hit rate, and one with poor hit rate. For these 3 results, there were 3 configuration pairs of cache size and block size that resulted in the data. Take the same 3 configuration pairs, but this time run the simulation with a fully associate cache, using LRU replacement policy. Compare the results obtained: The Direct Mapped good result versus the Fully Associative good result, the Direct Mapped medium result versus the Fully Associative medium result, and the Direct Mapped poor result versus the Fully Associative poor result. How much difference did the change to fully associative architecture make? Now change the replacement policy to Random replacement, and run the 3 tests again (using the same 3 configuration pairs). Does replacement policy make a significant difference? Record these 9 values in a new graph, with 3 lines: for Direct Mapped, for Fully Associative-LRU and for Fully Associative-Random. Note that this step is only for **column-major averaging** so that by the end of the lab you'll have  $(2+1) \times 2 = 6$  graphs.
- c) **N-way Set Associative Caches:** To save on hardware costs, fully set-associative caches are rarely used. Instead, most of the benefit can be obtained with an N-way set associative cache. Pick the medium hit rate configuration that you found in a) and used again in b), and change the architecture to N-way set associative. For several different set sizes (at least 4) and LRU replacement policy, run the program and record the hit rate, miss rate and number of misses. What set size gives the best result? How much improvement is gained as N (the number of blocks in a set) increases each step? Now repeat the tests, but for the good hit rate configuration from a) and b). Record these data and answer the same question again. Finally, repeat for the poor hit rate configuration.



**Figure 8.18** Miss rate versus block size and cache size on SPEC92 benchmark

Adapted with permission from Hennessy and Patterson, *Computer Architecture: A Quantitative Approach*, 5th ed., Morgan Kaufmann, 2012.